

# Understanding, Cleaning and Visualizing Data Through Pandas

## Data Cleaning and Preprocessing

### 1. Detailed Parsing of Salary Information:

- **Remove Currency Symbols:** Remove "\$" or other currency symbols from "Salary Estimate" to simplify the salary parsing process.
- **Salary Type Extraction:** Create a column that identifies whether the salary is presented as "per year," "per hour," etc., based on common indicators in "Salary Estimate."

### 2. Handle Special Characters in Text Columns:

- **Text Standardization:** Remove or replace special characters (e.g., "\n", "\t", and other symbols) from columns like "Job Description" and "Company Name" for consistency.
- **Lowercasing:** Convert all text columns (e.g., "Job Title," "Company Name," "Industry") to lowercase to standardize the format and avoid case sensitivity issues during analysis.

### 3. Company Name Refinement:

- **Remove Rating from Company Name:** Clean up "Company Name" by removing any appended rating values (e.g., "Company Name\n3.8").
- **Unique Company Identifier:** Generate a unique identifier for each company based on the cleaned "Company Name" to help with comparisons.

### 4. Standardizing Location Data:

- **Expand State Abbreviations:** For the "Location" column, expand state abbreviations (e.g., "CA" → "California") to ensure uniformity.
- **Location Format Consistency:** Check for any inconsistent formats in "Location" and "Headquarters" and standardize to "City, State" format.
- **Distance Calculation:** If city and state information are available, create approximate distance calculations between "Location" and "Headquarters" to identify remote positions.

### 5. Revenue Column Parsing and Cleaning:

- **Extract Revenue Range:** For each value in "Revenue," extract minimum and maximum revenue values to separate columns.
- **Currency Standardization:** Ensure all revenue values are in the same currency (e.g., USD) and handle conversions if any foreign currencies are detected.
- **Fill Missing Values in Revenue:** Fill missing values in the "Revenue" column.

## 6. Missing Data Handling:

- **Fill Missing Values in Categorical Columns:** Fill missing values in columns such as "Industry" and "Sector."
- **Fill Missing Values in Rating:** Fill missing values in "Rating."

## 7. Cleaning and Extracting Data from "Job Description":

- **Length Calculation:** Create a column to calculate the word count or character length of each "Job Description" to help filter lengthy descriptions.
- **Remote Work Flag:** Identify if the job allows remote work by checking for keywords (like "remote," "telecommute") in the "Job Description".
- **Experience Level Parsing:** Add columns to identify required experience levels (e.g., "entry level," "5+ years") by parsing "Job Description."

## 8. Additional Column Standardization:

- **Numeric Conversion of Founding Year:** Ensure "Founded" is a numeric column, and set a minimum year threshold (e.g., 1800) to filter out any erroneous entries.
- **Size Categorization:** For "Size," create size categories (e.g., "small," "medium," "large") based on ranges in employee count.

## 9. Outlier Detection and Handling:

- **Salary Outliers:** Detect outliers in the salary range and flag unusually high or low salary estimates for each job title.
- **Rating Outliers:** Identify extreme ratings (e.g., below 2.0 or above 4.5) and flag them to verify their accuracy.
- **Founded Year Outliers:** Flag companies with founding years outside a realistic range as potential data entry errors.

## 10. Column Optimization and Data Type Conversion:

- **Convert Numerical Columns:** Ensure columns like "Rating" and "Founded" are in a suitable numeric format (`int64` or `float64`) and convert others as appropriate.
- **Memory Optimization:** Downcast numeric columns (e.g., "Rating," "Founded") to the smallest integer or float type without data loss to optimize memory usage.

---

## Exploratory Data Analysis (EDA)

### 1. Salary and Rating Analysis:

- **Salary Range Analysis:** Calculate the median, mean, and standard deviation of salary ranges across different job titles, industries, and sectors.
- **Top and Bottom Companies by Salary:** Identify companies that offer the highest and lowest average salary ranges.
- **Rating Distribution Analysis:** Calculate the proportion of companies with ratings in different ranges (e.g.,  $<3.0$ ,  $3.0-4.0$ ,  $>4.0$ ).

## 2. Sector and Industry Analysis:

- **Sector Diversity:** Count the number of unique industries within each sector, and analyze if certain sectors have a higher concentration of industries.
- **Top Industries for Data Science:** List the top 5 industries with the most job postings for data science roles and their average salaries and ratings.
- **Revenue by Sector and Industry:** Calculate the average revenue for each sector and industry to identify the most profitable areas.

## 3. Location and Company Insights:

- **Top Locations by Job Count:** Identify the top 10 locations (cities and states) with the highest job counts and analyze their average salaries and ratings.
- **Headquarters vs. Job Location:** Compare ratings for jobs at a company's headquarters vs. remote offices.
- **Company Size and Location Correlation:** Explore whether certain locations are more likely to host larger or smaller companies.

## 4. Experience Level Analysis:

- **Job Title Levels:** Calculate the average rating and salary for each level (e.g., "Junior," "Mid-level," "Senior").
- **Experience Level and Location:** Analyze if certain experience levels (e.g., entry-level, senior) are more commonly available in specific locations or industries.
- **Experience Level by Sector:** Analyze the distribution of experience levels by sector and identify if certain sectors focus more on senior or junior roles.

## 5. Tricky EDA:

- **Founding Year vs. Salary:** Analyze if older companies tend to offer higher or lower salaries compared to newer companies.
- **Remote Work and Rating:** Compare the ratings of jobs marked as remote vs. non-remote to see if remote positions correlate with higher ratings.
- **Competitors and Salary:** Check if companies with multiple listed competitors tend to offer higher salaries on average.

---

## Data Transformation

### 1. Advanced Keyword Tagging:

- **Add Skill Flags:** Create binary columns to flag if certain skills (e.g., "Python," "SQL," "machine learning") are mentioned in the "Job Description."
- **Identify Benefits:** Add columns for common benefits (e.g., "health-care," "retirement plan") based on keywords found in the "Job Description."
- **Remote Job Tagging:** Add a column to indicate if the job allows remote work based on keywords in "Job Description" or "Location."

### 2. Time-Based Transformation:

- **Convert Founded Year to Age:** Calculate each company's age by subtracting the founding year from the current year.
- **Establish Company Age Categories:** Categorize companies into age groups (e.g., "New" for <10 years, "Established" for 10-50 years, "Legacy" for >50 years).
- **Estimate Job Posting Age:** If date information is available, calculate how long ago each job posting was listed.

### 3. Categorization and Grouping:

- **Sector and Industry Mapping:** Create a mapping dictionary to reclassify "Industry" values into broader "Sector" categories if needed.
- **Location Grouping:** Group locations into regions (e.g., Northeast, Midwest) for U.S. data or other relevant geographical groupings.
- **Company Size Bucketing:** Create buckets for company size (e.g., "Small" for <500 employees, "Medium" for 500-5000, "Large" for >5000) and add a new column.

### 4. Advanced Revenue and Salary Conversion:

- **Calculate Revenue per Employee:** Create a column for revenue per employee by dividing the revenue by the employee size estimate.
- **Salary Normalization:** Normalize salary ranges (min and max) to a single yearly estimate for comparison, especially for hourly roles.
- **Create Salary Range Buckets:** Classify salary ranges into brackets (e.g., "<\$50k," "\$50k-\$100k," ">\$100k") to enable clearer comparisons.

### 5. Competitor Data Transformation:

- **Competitor Count:** Add a column indicating the number of competitors each company lists.

- **Frequent Competitors:** Add a flag for companies that list commonly seen competitors and analyze if these companies have different rating or salary trends.
- **Common Competitor Networks:** Create clusters for companies with the same competitors, allowing analysis of companies within specific competitive networks.

---

## Comparative Analysis

### 1. Sector vs. Industry Comparison:

- **Rating Comparison:** Compare the average rating across industries within each sector to see if certain industries outperform others within the same sector.
- **Salary Range by Industry:** Compare the median salary ranges for each industry within top sectors.
- **Revenue vs. Sector:** Analyze if certain sectors tend to have higher revenue brackets compared to others.

### 2. Ownership Type Analysis:

- **Public vs. Private Salary Comparison:** Compare average salaries between public and private companies.
- **Rating Distribution by Ownership:** Compare rating distributions across different ownership types to see if public or nonprofit organizations have higher ratings.
- **Company Size by Ownership:** Analyze if public or private companies tend to be larger or smaller in terms of employee size.

### 3. Company Age vs. Rating:

- **Compare Ratings by Age Group:** Analyze average ratings of companies in each age group to see if older companies have higher ratings.
- **Rating Stability Over Time:** Examine if companies founded before 2000 have more consistent (less variable) ratings than newer companies.
- **Sector Popularity by Age:** Compare the prevalence of different sectors within each age group, observing which sectors are popular among older vs. newer companies.

### 4. Location-Based Comparison:

- **Salary by Region:** Compare average salary ranges across regions to see if certain regions offer higher pay for data science jobs.
- **Remote vs. Onsite Roles:** Compare salary ranges, job counts, and ratings for remote vs. onsite positions.



- **Job Count by State:** Compare the total job count and average salary for each state to identify states with the highest job availability.

#### 5. Tricky Comparative Analysis:

- **Competitor Impact on Salary:** Compare salary ranges for companies with and without listed competitors to analyze if competition impacts salaries.
- **Industry Longevity Analysis:** Compare average founding years by industry to see if certain industries have more established companies.
- **Rating and Salary Disparity:** Analyze if companies with higher ratings offer salaries that deviate more from industry or sector norms.

### Advanced Data Insights and Parsing

#### 1. Competitor Relationships and Analysis:

- Identify and count the most frequently listed competitors across all companies. Create a column to show how many times each competitor is mentioned.
- For companies with multiple competitors listed, calculate the average rating difference between them and their competitors. Are highly-rated companies more likely to list high-profile competitors?
- Create a column indicating whether a company shares competitors with any top-rated companies (Rating 4), and analyze if this has any effect on its own rating.

#### 2. Industry and Sector Dynamics:

- Determine the sectors with the most diversity in industries and analyze if this correlates with higher salary estimates or job counts.
- Find the industries with the highest average ratings and analyze which specific sectors they belong to. Does belonging to a specific sector influence the industry's average rating?
- Identify industries with the highest number of job postings and check if these align with high salary estimates. Is there a trend between industry popularity and salary?

#### 3. Company Size and Growth Patterns:

- Group companies by "Size" and analyze the average founding year for each size group. Are larger companies generally older, or is there a trend of newer, fast-growing companies in the dataset?
- Create an "Employee Growth Index" by dividing the number of employees (from "Size") by the number of years since founding. Compare this index across sectors to identify fast-growing sectors.

- For each company, calculate the ratio of the company size to the number of competitors listed. Does a larger company tend to have more competitors?

#### 4. Revenue and Financial Analysis:

- For companies with available revenue data, calculate the revenue per employee (using "Revenue" and "Size") and analyze which sectors have the highest revenue per employee.
- Determine the average revenue for each sector and examine if high-revenue sectors have higher-rated companies on average.
- Identify companies with the widest revenue ranges and examine if these companies have more job postings or higher ratings.

#### 5. Salary Insights and Trends:

- For each job title, calculate the average salary range width (difference between max and min salary) to see which roles have the most salary variability.
- Identify the top 5 job titles with the highest average salary ranges and analyze if they are concentrated in certain sectors or locations.
- Compare median salary ranges between private and public companies to determine if there's a significant difference based on ownership type.

#### 6. Job Description Keywords and Popularity:

- Identify the top 10 keywords in "Job Description" across different sectors. Are certain keywords unique to specific sectors?
- Parse common industry-specific terms from "Job Description" (e.g., "fintech" for finance, "sustainability" for environmental) and analyze if jobs mentioning these terms have higher average ratings.
- Check the percentage of job descriptions mentioning benefits like "remote," "flexible hours," or "bonus," and analyze if jobs mentioning these perks have higher average ratings.

#### 7. Competitor and Rating Correlation:

- Calculate the average rating for companies that list specific high-profile competitors (e.g., Google, Microsoft) and compare it to the ratings of companies without those competitors.
- Create a "Competitor Count Impact Score" by examining if companies with a higher number of competitors listed tend to have higher or lower ratings.
- Analyze if companies with unique competitors (listed only once) have distinct characteristics in ratings or salary compared to those with common competitors.

## Visualization

### 1. Basic Visualizations:

- Plot the distribution of company sizes (in terms of employee numbers) and see how it aligns with "Type of Ownership" (e.g., private, public, nonprofit).
- Create a bar chart showing the average rating for each "Type of Ownership" to see if ownership types correlate with employee satisfaction.

### 2. Location-Based Visualizations:

- Generate a bar chart to compare the number of job postings across the top 10 cities, highlighting popular cities for data science roles.
- Create a map plot to show the geographic distribution of job postings by state, using color intensity to represent job density.

### 3. Salary and Rating Analysis:

- Create a scatter plot showing the relationship between average salary range and company ratings, and add a trend line to analyze if higher-rated companies tend to offer higher salaries.
- Plot a histogram of salary ranges for different sectors, allowing comparison of sector-based salary distributions.
- Generate a heatmap to visualize the correlation matrix between numerical columns (e.g., ratings, salary estimates, founding year) to identify potential relationships.

### 4. Tricky Visualizations:

- Plot a line graph of the average rating of companies over time by "Founded Year" to see if there's a trend in ratings for companies founded in different decades.
- Create a box plot showing the salary range distributions by industry, highlighting which industries have the widest salary range.
- Generate a stacked bar chart of job titles within each sector to observe which titles are most common in different industries.

### 5. Competitor and Popularity Insights:

- Create a network graph to show connections between companies and their competitors, with node sizes based on the company's average rating.
- Plot a bar chart showing the top 10 most common competitors listed, along with the average rating of companies that have these competitors.

### 6. Company and Industry Focused Visualizations:

- Create a stacked area chart to show how job postings by company size ("Size") are distributed over time (based on "Founded Year").



- Plot a donut chart representing the distribution of job postings by "Type of Ownership" to visualize how job availability varies by ownership type.
  - Generate a radar chart (spider chart) comparing key metrics (e.g., average rating, median salary, company size) across top sectors to easily compare sector strengths.
-